

Methodology: GPU DNS of Fractal-Interface Mixing in a Developing Duct

TorChannel — MERGE fractal-mixing study

June 30, 2026

Abstract

We describe the numerical methodology used to test whether a fractal (Koch) scalar interface enhances mixing in a laminar developing pipe flow, and how the Schmidt number Sc is swept to probe the high-Peclet regime. The pipeline couples a GPU-accelerated, double-precision direct numerical simulation (DNS) of the incompressible Navier–Stokes equations — staggered grid, fractional-step pressure projection, IMEX time integration, Brinkman-penalized immersed pipe wall — to a conservative passive-scalar transport solver. To make the high- Sc , high-fractal-generation sweep tractable we exploit three physical properties of this configuration (parallel fully-developed flow, high-Peclet parabolicity, and an N -independent wall) that collapse the cost from $O(Pe)$ pseudo-time steps per case to a single downstream sweep on a frozen, once-computed velocity field.

1 Problem setup

The configuration is a short, developing, laminar duct of square outer cross-section $L_y \times L_z = 1 \times 1$ and streamwise length L_x , inside which a circular pipe of radius $R = 0.42$ is imposed as an immersed solid. Two co-flowing streams of a passive scalar $c \in [0, 1]$ enter the pipe at the inlet, separated by an interface. In the *baffle* configuration the dividing interface is a *Koch fractal* of generation N prescribed at the inlet; the pipe wall itself remains smooth and circular. The question is whether increasing the interfacial generation N — which multiplies the interfacial *area* without changing the 50/50 volume split — accelerates the downstream homogenisation of c , and how that effect scales with the Schmidt number

$$Sc = \frac{\nu}{D}, \quad Pe = Re \, Sc = \frac{U_b L}{D}, \quad (1)$$

where ν is the kinematic viscosity, D the scalar diffusivity, U_b the bulk velocity and L a reference length. The molecular regime of interest (aqueous solutes) is $Sc \sim 10^3$, i.e. Pe up to $\sim 10^4$ at the chosen $Re = 40$.

2 Governing equations

The flow obeys the incompressible Navier–Stokes equations, with the passive scalar advected by the resulting divergence-free velocity:

$$\nabla \cdot \mathbf{u} = 0, \quad (2)$$

$$\frac{\partial \mathbf{u}}{\partial t} + (\mathbf{u} \cdot \nabla) \mathbf{u} = -\nabla p + \nu \nabla^2 \mathbf{u} + \mathbf{f}_{\text{ib}}, \quad (3)$$

$$\frac{\partial c}{\partial t} + \mathbf{u} \cdot \nabla c = D \nabla^2 c, \quad D = \frac{\nu}{Sc}. \quad (4)$$




Figure 1: **(Sketch, schematic.)** Configuration. Two co-flowing scalar streams enter an immersed circular pipe separated by a Koch-folded interface (inset, cross-section). The fully-developed velocity is frozen and the steady scalar is obtained by a single downstream parabolic sweep.

All fields are stored in double precision throughout. The scalar is *passive*: it does not feed back into the momentum equation, so (4) carries no projection step. $\mathbf{f}_{\text{ib}}$ is the immersed-boundary penalization force (Section 5). Because the conservative flux form $\nabla \cdot (\mathbf{u}c)$ equals $\mathbf{u} \cdot \nabla c$ for a divergence-free field, advection is discretised in flux form to conserve $\langle c \rangle$ exactly.

3 Spatial discretisation

Staggered MAC grid. Velocities live on cell faces (the streamwise u on x -faces, v on y -faces, w on z -faces) while pressure and the scalar c are cell-centred. This Marker-and-Cell arrangement avoids odd–even pressure decoupling and gives a compact, exactly conservative divergence/gradient pair. All fields carry a single layer of ghost cells (indices 0 and -1 in each direction) through which boundary conditions are imposed.

Grid. The mesh is uniform in the periodic-capable x and y directions ($dx = L_x/n_x$, $dy = L_y/n_y$) and *stretched* in the wall-normal z via a hyperbolic-tangent mapping, clustering points near the walls; the variable spacings dz_c (centre-to-centre) and dz_f (face-to-face) enter every z -derivative. Spatial operators are second-order: central differences for diffusion, and conservative flux differences for advection. Performance-critical kernels are fused for the GPU.

4 Time integration: fractional-step projection

Momentum is advanced by a semi-implicit (IMEX) fractional-step (projection) method. Within one step of size Δt :

1. **Explicit terms (Adams–Bashforth 2).** Advection and the in-plane (x, y) diffusion are evaluated explicitly and combined with second-order Adams–Bashforth,

$$\mathbf{u}^* = \mathbf{u}^n + \Delta t \left(\frac{3}{2} \mathbf{R}^n - \frac{1}{2} \mathbf{R}^{n-1} \right), \quad \mathbf{R} = \nu \nabla_{xy}^2 \mathbf{u} - (\mathbf{u} \cdot \nabla) \mathbf{u} + \mathbf{f}, \quad (5)$$

(the first step falls back to forward Euler).

2. **Implicit wall-normal diffusion (θ -method).** The stiff z -diffusion is treated implicitly, $(I - \theta \Delta t \nu \partial_{zz}) \mathbf{u}^{**} = \mathbf{u}^* + (1 - \theta) \Delta t \nu \partial_{zz} \mathbf{u}^*$, with $\theta = \frac{1}{2}$ (Crank–Nicolson). On the stretched grid this is a tridiagonal system per wall-normal column, solved by a batched Thomas solver over all columns at once. Treating only the z -diffusion implicitly removes the most restrictive viscous stability limit while keeping the linear solve one-dimensional.

3. **Immersed-boundary penalization** (Section 5) is applied to $\mathbf{u}^{**}$ before projection.

4. **Pressure projection.** A Poisson equation enforces incompressibility,

$$\nabla^2 p = \frac{1}{\Delta t} \nabla \cdot \mathbf{u}^{**}, \quad \mathbf{u}^{n+1} = \mathbf{u}^{**} - \Delta t \nabla p, \quad (6)$$

solved by an FFT-based fast Poisson solver in the periodic x, y planes combined with a tridiagonal solve in z (a direct banded solver is available as a fallback). The resulting $\mathbf{u}^{n+1}$ is discretely divergence-free.

5 Immersed pipe wall (Brinkman penalization)

The circular pipe is embedded in the Cartesian box by an implicit Brinkman (volume-penalization) force that drives the velocity to zero inside the solid:

$$\mathbf{f}_{\text{ib}} = -\frac{\chi}{\eta} \mathbf{u}, \quad \chi(\mathbf{x}) = \begin{cases} 1 & \text{solid (outside the pipe)} \\ 0 & \text{fluid (inside the pipe)} \end{cases} \quad (7)$$

with a small permeability $\eta = 10^{-4}$. The penalization is applied implicitly to the intermediate velocity *before* the projection, $\mathbf{u} \leftarrow \mathbf{u}/(1 + \Delta t \chi/\eta)$, so the fast FFT-Poisson solver is left untouched. The solid mask χ is sampled separately on each staggered face (χ_u, χ_v, χ_w) and at cell centres (χ_c); here it is the complement of the inscribed disc of radius $R = 0.42$ centred on the cross-section, giving a solid fraction ≈ 0.45 . The bulk-flux controller and all mixing diagnostics are restricted to fluid cells ($\chi_c < \frac{1}{2}$) so that scalar diffusing into the (non-physical) solid is never counted as mixing.

6 Inflow / outflow boundary conditions

The duct is open (`bc_x = inout`): periodic in y only where specified, walls otherwise, and a prescribed inlet / convective outlet in x . The inlet streamwise velocity is a smooth duct profile, the product of parabolae in y and z that vanish at the walls,

$$u_{\text{in}}(y, z) \propto \frac{4y(L_y - y)}{L_y^2} \frac{4z(L_z - z)}{L_z^2}, \quad (8)$$

masked to zero inside the solid and normalised so the fluid-area-mean inlet velocity equals U_b . The outflow imposes zero streamwise gradient with a global mass-flux correction enforcing $\int u_{\text{out}} = \int u_{\text{in}}$ over the fluid face, so the discrete domain conserves mass. The scalar inlet ghost is pinned to the prescribed interface $c_{\text{inlet}}(y, z)$ each step; the scalar outflow is zero-gradient. Lateral walls use no-flux ($\partial_n c = 0$) for the scalar.

7 The baffle: a Koch-fractal inlet interface

In the *baffle* mode the pipe wall is smooth; the only N -dependent ingredient is the *scalar interface* prescribed at the inlet. The two streams meet on a curve that runs wall-to-wall across the pipe and is Koch-folded in the y -direction at generation N .

Generator. The interface is built from an area-balanced four-segment zigzag motif, each segment of length $1/r$:

$$(0, 0) \rightarrow (\tfrac{1}{4}, a) \rightarrow (\tfrac{1}{2}, 0) \rightarrow (\tfrac{3}{4}, -a) \rightarrow (1, 0), \quad a = \sqrt{(1/r)^2 - (1/4)^2},$$

valid for $1 < r < 4$. Each outward triangle is matched by an equal inward indentation, so the fold preserves the exact 50/50 area split at every generation (the mean scalar is held at $\langle c \rangle = \frac{1}{2}$). Replacing every edge by an affine copy of the motif N times yields a generation- N polyline with fractal dimension

$$D_f = \frac{\log 4}{\log r}. \quad (9)$$

The interfacial length (hence the area available for diffusive flux) grows as $\sim (4/r)^N$ relative to the flat $N = 0$ baseline, while the enclosed area is unchanged — isolating the effect of interfacial area on mixing.




Figure 2: Area-balanced Koch generator ($r = 3$, $D_f = \log 4 / \log 3 \approx 1.26$). Each edge is replaced by the four-segment motif, multiplying interfacial length by $\sim 4/r$ per generation while preserving the 50/50 split. The inlet interface uses generation N of this fold; these are the actual polylines imposed.

Rasterisation. The cross-sectional field $c(y, z)$ is set from the signed distance ϕ to the open Koch polyline, smoothed over a prescribed interface thickness ε (a fixed number of cells):

$$c(y, z) = \frac{1}{2} [1 + \tanh(\phi/\varepsilon)]. \quad (10)$$

The distance *magnitude* comes from the open polyline, but its *sign* (which stream a cell belongs to) is taken from an even-odd point-in-polygon test against the interface closed along the far wall. This is robust at high N , where a naive nearest-segment side test produces wrong-sign corner artefacts ($N \geq 3$). The interface is homogeneous in x at the inlet; all streamwise structure that develops downstream is produced by the flow. A companion *surface_baffle* mode additionally folds the *pipe wall* into a Koch surface, differing from *baffle* only by the wall, for a controlled isolation of the surface-roughness effect.

8 Passive-scalar transport

The scalar lives at cell centres and is advanced with the same IMEX split as the momentum equations (AB2 for advection and in-plane diffusion; θ -method for wall-normal diffusion), reusing the batched tridiagonal solver. Advection is in conservative flux form,

$$\nabla \cdot (\mathbf{u}c) = \frac{F_{i+\frac{1}{2}}^x - F_{i-\frac{1}{2}}^x}{dx} + \frac{F_{j+\frac{1}{2}}^y - F_{j-\frac{1}{2}}^y}{dy} + \frac{F_{k+\frac{1}{2}}^z - F_{k-\frac{1}{2}}^z}{dz_f}, \quad (11)$$

with face fluxes $F = U_{\text{face}} c_{\text{face}}$ built from the MAC face velocities. Two face reconstructions are available:

- **Central** (c_{face} = linear average) — second-order, used at moderate cell-Peclet.
- **van Leer TVD** — a MUSCL upwind reconstruction with the van Leer limiter, $c_{\text{face}} = c_{\text{up}} + \frac{1}{2} \psi$, monotone (no dispersive over/undershoot) at the high cell-Peclet of high- Sc runs. Since the face velocity is identically zero at every wall, the limiter never alters a wall flux.

The wall-normal boundary condition is no-flux (Neumann, $\partial_z c = 0$), which conserves the scalar mean and lets the segregation variance decay to zero — the correct closure for a mixing / decay study. A Dirichlet ($c = 0$) option exists for future sustained-gradient setups.

9 Mixing diagnostics

Mixing is quantified by the *intensity of segregation*, the fluid-volume-weighted, variance-based mixedness

$$M = \frac{\sigma_c}{\sigma_{\max}}, \quad \sigma_c^2 = \langle (c - \bar{c})^2 \rangle, \quad \sigma_{\max} = \sqrt{\bar{c}(1 - \bar{c})}, \quad (12)$$

where $\langle \cdot \rangle$ is the z -stretch-weighted average over fluid cells only. $M = 1$ is fully segregated (two pure streams); $M \rightarrow 0$ is fully mixed. For the developing duct the relevant object is the per-slice profile $M(x)$, evaluated plane-by-plane, from which a *mixing length* L_{mix} is read off as the streamwise station where $M(x)$ first falls below a threshold (e.g. $M = 0.9$), linearly interpolated between planes. The fractal effect is reported as the ratio $L_{\text{mix}}(N)/L_{\text{mix}}(0)$. A complementary scalar-dissipation diagnostic $\langle |\nabla c|^2 \rangle$ is also available; unlike the gravest-mode-dominated variance, it is directly sensitive to interfacial area and so captures the fractal's contribution that the variance metric alone can miss.

10 Acceleration: from $O(Pe)$ to one sweep

A direct time-march of (4) to steady state costs $\sim Pe$ pseudo-time steps and, at $Sc = 10^3$ on a fine cross-section, is infeasible. Three properties of this configuration collapse the cost.

10.1 Parabolic space-marching (high-Peclet) scalar solver

In a developing duct with $u > 0$ everywhere (no streamwise recirculation), the steady scalar balance

$$u \frac{\partial c}{\partial x} = D \left(\frac{\partial^2 c}{\partial y^2} + \frac{\partial^2 c}{\partial z^2} \right) - \left(v \frac{\partial c}{\partial y} + w \frac{\partial c}{\partial z} \right) \quad (13)$$

is obtained after *dropping the streamwise diffusion* $\partial_{xx}c$ — the *parabolic* (boundary-layer) approximation, asymptotically exact as $Pe \rightarrow \infty$ and, unlike a time-march, free of streamwise *numerical* diffusion. The equation is then parabolic in x and the steady field follows from a *single downstream sweep* — one cross-plane solve per x -station — instead of marching the unsteady field to steady state. The streamwise term uses second-order upwind (BDF2) from the two already-solved upstream planes; each cross-plane (y, z) diffusion solve uses alternating-direction implicit (ADI) line relaxation (a z -implicit then a y -implicit tridiagonal sweep), which converges in a handful of inner iterations — fastest precisely at high Sc , where the streamwise term dominates the diagonal. Cost is $O(n_x)$ plane solves rather than $O(Pe)$ full-field updates. For the smooth baffle $v = w = 0$, so transverse advection is switched off entirely.

10.2 Frozen, fully-developed velocity (parallel-flow assumption)

The scalar marcher runs on a *frozen* velocity field. The pipe wall is straight and x -independent, so the laminar flow becomes *fully developed* a short distance past the inlet: beyond the entrance length the profile $u(y, z)$ stops changing with x ($\partial_x u \approx 0$) and every downstream plane is identical. We verified this directly on the computed field — the cross-sectional profile is within 0.1% of its developed shape by $x \approx 2.8$ and $\partial_x u \sim 5 \times 10^{-5}$ thereafter. Consequently the velocity need only be solved over the short developing region, and extending it to an arbitrarily long scalar domain is exact: the marcher edge-clamps the streamwise interpolation to the developed profile, i.e. it repeats the one converged $u(y, z)$ downstream. This is the same parallel-flow premise that justifies dropping $\partial_{xx}c$ in the first place; it holds because the baffle wall does not vary with x (it would *not* hold for a streamwise-varying cross-section).

10.3 N -independent velocity, solved coarse and reused

In the baffle mode the Koch generation N lives only in the inlet *scalar*; the wall — hence the entire velocity field — is identical for every N . The developed velocity is therefore solved *once* and cached, then reused across the whole N -sweep. Moreover the smooth-pipe flow is well resolved on a *coarse* cross-section, whereas only the scalar needs the fine grid to resolve the Koch striations without false diffusion. We exploit this by solving the velocity on a coarse mesh and tri-linearly interpolating it (zeroed inside the fine solid mask) onto the fine scalar grid. The full N -sweep thus costs *one* coarse velocity solve plus N cheap parabolic scalar sweeps, in place of one fine velocity solve and a high- Sc time-march *per* N .

11 The Schmidt-number sweep

A sweep fixes the geometry (R , L_x , grid) and the threshold, and varies the Koch generation $N = 0, 1, 2, \dots$ at a target Schmidt number. For each N the parabolic marcher returns the steady $c(x, y, z)$ on the frozen developed velocity; from it we extract $M(x)$, the mid-plane and cross-section slices, and $L_{\text{mix}}(N)$. Comparing $L_{\text{mix}}(N)/L_{\text{mix}}(0)$ across N , at increasing Sc (equivalently $Pe = Re Sc$), tests whether the fractal interface’s extra area shortens the mixing length and whether that advantage persists into the molecular ($Sc \sim 10^3$) regime.

Validation. The scalar solver is checked against an analytic error-function diffusion profile, exact mean conservation, and zero numerical streamwise diffusion; the DNS divergence is monitored every step and the run aborts if it grows.

12 Preliminary results ($Sc = 10, 100$ and 1000)

Table 1 and Figure 3 collect the mixing length L_{mix} (the station where $M(x)$ first crosses the threshold) for the baffle, at the three Schmidt numbers run. The $Sc = 100$ and $Sc = 1000$ data come from the parabolic marcher; the $Sc = 10$ data from a direct time-march to steady state (an independent cross-check of the marcher). Because L_{mix} grows with Pe , the higher- Sc sweeps use longer domains ($L_x = 6, 72, 90$ for $Sc = 10, 100, 1000$) so that all N cross $M = 0.9$ within them; the threshold is therefore the *same* $M = 0.9$ for every Sc , an apples-to-apples comparison. The ratio $L_{\text{mix}}(N)/L_{\text{mix}}(0)$ *does* depend on the threshold — it rises smoothly toward 1 as M is lowered (Figure 4) — so every comparison here is made at a *fixed* threshold; at fixed M the N -ordering and trend are robust, and only the reported magnitude shifts with the choice of M .

N	$Sc = 10$ ($M=0.9$)		$Sc = 100$ ($M=0.9$)		$Sc = 1000$ ($M=0.9$)	
	L_{mix}	ratio	L_{mix}	ratio	L_{mix}	ratio
0	0.492	1.000	5.490	1.000	62.818	1.000
1	0.303	0.616	3.811	0.694	38.370	0.611
2	0.198	0.402	2.385	0.434	26.258	0.418
3	0.152	0.309	1.998	0.364	22.164	0.353
4	0.136	0.276	1.936	0.353	21.657	0.345
5	—	—	—	—	21.581	0.344

Table 1: Mixing length L_{mix} (station where $M(x)$ crosses the threshold) and its ratio to the flat $N = 0$ baseline, baffle mode, $R = 0.42$. All at the same threshold $M = 0.9$. The domain is sized so each field decays well past the threshold: $Sc = 10$ uses $L_x = 6$, $Sc = 100$ uses $L_x = 72$, $Sc = 1000$ uses $L_x = 90$ (longer domains for the higher Pe).




Figure 3: Fractal mixing enhancement, all at a fixed threshold $M = 0.9$: the mixing length falls with generation N and saturates near 0.35 by $N \approx 3$, *roughly independently of Sc* (the three curves nearly overlap; $Sc = 100$ is marginally the weakest). The grey dash-dotted line is the proposal’s diffusive-limit prediction $r^{-D_f N} = 4^{-N}$ (Eq. 4): the measured ratios fall well short of it — the finite- Pe duct realises only part of the ideal exponential cut. The ratio’s magnitude depends on the threshold (Figure 4), but its N -ordering does not.

The mechanism is visible directly in the cross-sections (Figure 6): higher N injects more interfacial length into the same disc, and the steep transverse scalar gradients it creates are ground down by diffusion over a shorter streamwise distance. The companion $M(x)$ decay (Figure 8) shows the same ordering as a single scalar per station, for both $Sc = 100$ and $Sc = 1000$.

13 Implications

1. **The fractal interface robustly shortens the mixing length, roughly independently of Sc .** At a fixed threshold ($M = 0.9$), folding to $N = 3$ cuts L_{mix} to 0.31, 0.36 and 0.35 of the flat baseline at $Sc = 10, 100, 1000$ — a $\sim 3\times$ reduction that saturates near 0.35 for $N \geq 3$ at every Sc , with the same volume split and the same wall. This confirms, in a resolved laminar duct, the core MERGE premise that interfacial *area* (not bulk stirring) drives the enhancement. (At a *common* threshold there is no clear monotone Sc -trend — the apparent strengthening with Sc in earlier mixed-threshold readings was a threshold artifact, see Fig. 4.)
2. **Diminishing returns.** Most of the gain is in the first folds ($N: 0 \rightarrow 1 \rightarrow 2$); the ratio then flattens, saturating by $N \approx 3-4$ at every Sc (e.g. $Sc = 1000$ gives 0.42, 0.35, 0.345, 0.344 at $N = 2, 3, 4, 5$). Folds finer than the local diffusion length are smeared away before they can laminate, so deeper fractals stop paying off.
3. **L_{mix} scales steeply with Pe .** The $N = 0$ baseline grows $0.49 \rightarrow 5.5 \rightarrow 63$ for $Sc = 10 \rightarrow 100 \rightarrow 1000$ (all at $M = 0.9$) — about a decade per decade of Sc . This is why a high- Sc run needs a long domain, and is the direct motivation for the parabolic-marcher / frozen-velocity acceleration of Section 10.
4. **Interfacial folding dominates over wall roughness.** The *surface_baffle* (Koch folds added to the pipe *wall*) gives a weaker reduction than the interface baffle at equal N (e.g. at $Sc = 10, N = 3$: $L_{\text{mix}} = 0.21$ vs. 0.15), indicating that the dominant mechanism is




Figure 4: Threshold sensitivity of the ratio $L_{\text{mix}}(N)/L_{\text{mix}}(0)$: each curve is a generation N (colour as in Fig. 8) swept over the mixedness threshold M . The ratio rises smoothly and *monotonically* toward 1 as M is lowered — it never flickers or reverses — and the N -ordering is preserved at every M . The fractal advantage is therefore largest near the inlet (high M) and fades toward complete mixing, but the qualitative conclusions are threshold-independent; only the reported magnitude depends on the chosen M . Over the same range the $Sc = 10$ and $Sc = 100$ panels have nearly the same shape, confirming the behaviour is not Sc -specific (see also $Sc = 1000$, Fig. 5).

lamination of the scalar interface, not corrugation-driven secondary flow.

5. **The acceleration is what makes the proposal regime reachable.** Because the velocity is fully developed and N -independent, the entire multi- N , high- Sc sweep collapses to one coarse velocity solve plus cheap parabolic scalar sweeps. Without it, true $Sc \sim 10^3$ by direct time-march would be infeasible; with it, the $Sc = 10^3$ sweep runs in minutes per generation.

Caveats. All ratios are quoted at a common $M = 0.9$. The $Sc = 10$ values come from a time-march on a coarse grid (some cases stopped at the step cap, not fully steady) while $Sc = 100$ and $Sc = 1000$ use the parabolic marcher on fine grids, so cross- Sc comparisons are indicative rather than exact — in particular the $Sc = 10$ ratios sit slightly below the marcher cases and should not be over-read as a trend. The configuration and generator diagrams (Figs. 1, 2) are schematic; the quantitative figures (Figs. 3, 4, 6, 8) are computed results.




Figure 5: Threshold sensitivity at $Sc = 1000$ ($N = 1$ – 5 , same quantity as Fig. 4). The reachable band is narrow ($M \gtrsim 0.88$) because the $L_x = 90$ field mixes $N = 0$ only to $M \approx 0.88$; within it the ratio is again smooth and monotone with the same N -ordering, consistent with the lower- Sc panels.

$Sc = 1000$ baffle: $c(y, z)$




Figure 6: Steady scalar $c(y, z)$ in the pipe cross-section at $Sc = 1000$, from the parabolic marcher. Rows: Koch generation $N = 0$ –5 (interfacial folding increases downward). Columns: streamwise station $x = 0$ (inlet) to $x = 90$ (outlet). The inlet interface ($x = 0$) is the prescribed generation- N Koch fold; downstream the fractal striations diffuse and the field homogenises faster at higher N (the per-panel M confirms the trend). Same 50/50 split and smooth wall for every N .

Sc= 100 baffle: $c(y, z)$




Figure 7: As Fig. 6 but at $Sc = 100$ ($N = 0$ – 4 , stations $x = 0$ to 72). At the lower Pe the field mixes further within the domain (outlet $M \approx 0.54$ – 0.57); the same trend holds — the fractal striations diffuse out and higher N homogenises sooner.




Figure 8: Segregation decay $M(x)$, generations N overlaid, for $Sc = 100$ ($L_x = 72$, left) and $Sc = 1000$ ($L_x = 90$, right). Higher N decays faster and further in both. Both cross $M = 0.9$ within their domains, so the mixing length is read at the same threshold for every Sc (Table 1).